

SHANGHAI JIAO TONG UNIVERSITY

CS353 LINUX KERNEL

Project 2B: Process Management

Chao Gao

5142029014

April 5, 2017

1 Introduction

Recently we have learned something about process management. Now I have a brief understanding about linux process descriptor, process ID and how processes are organized.

Let's recap the requirement of this project.

Process: schedule in times.

Add ctx, a new member to task_struct to record the schedule in times of the process. When a task is scheduled in to run on a CPU.

Export ctx under /proc/XXX/ctx.

2 Environment

Virtual OS: Ubuntu 16.04.

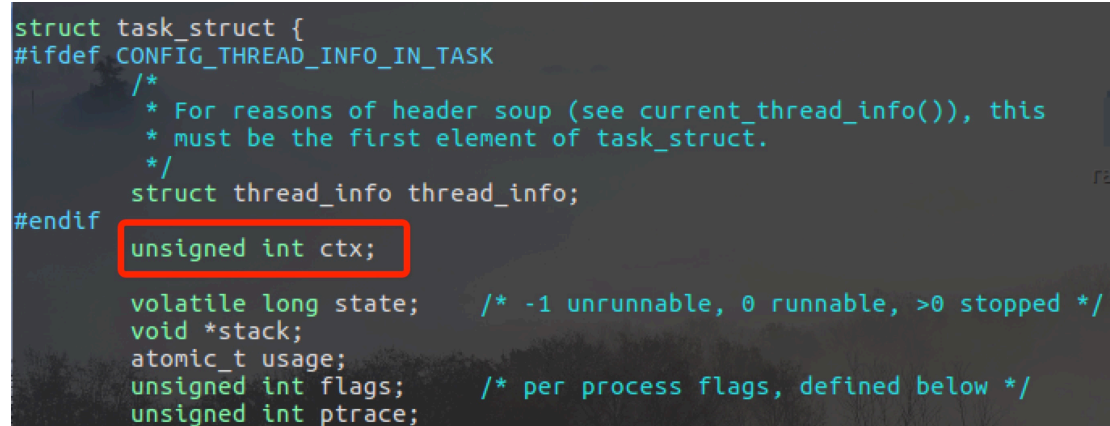
kernel version: 4.9.13.

3 Overall procedure

Here I will present the whole process of this project.

3.1 add ctx to task_struct

task_struct is in include/linux/sched.h.



```
struct task_struct {
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /*
     * For reasons of header soup (see current_thread_info()), this
     * must be the first element of task_struct.
     */
    struct thread_info thread_info;
#endif
    unsigned int ctx;
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags; /* per process flags, defined below */
    unsigned int ptrace;
```

3.2 ctx initialization

Initialize ctx=0 in function `_do_fork` in `kernel/fork.c`.

```
if (!IS_ERR(p)) {
    struct completion vfork;

    p->ctx = 0;

    struct pid *pid;

    trace_sched_process_fork(current, p);

    pid = get_task_pid(p, PIDTYPE_PID);
    nr = pid_vnr(pid);

    if (clone_flags & CLONE_PARENT_SETTID)
        put_user(nr, parent_tidptr);

    if (clone_flags & CLONE_VFORK) {
        p->vfork_done = &vfork;
        init_completion(&vfork);
        get_task_struct(p);
    }
}
```

3.3 plus ctx

When the processes are switched, plus ctx by 1 in `kernel/sched/core.c`.

```
if (likely(prev != next)) {
    rq->nr_switches++;
    rq->curr = next;
    ++*switch_count;

    next->ctx++;

    trace_sched_switch(preempt, prev, next);
    rq = context_switch(rq, prev, next, cookie); /* unlocks the rq */
} else {
    lockdep_unpin_lock(&rq->lock, cookie);
    raw_spin_unlock_irq(&rq->lock);
}

balance_callback(rq);
```

3.4 export ctx

Define a function `proc_pid_ctx` in `/fs/proc/base.c`

```
static const struct dentry_operations tid_map_files_dentry_operations = {
    .d_revalidate = map_files_d_revalidate,
    .d_delete     = pid_delete_dentry,
};

static int proc_pid_ctx(struct seq_file *m, struct pid_namespace *ns, struct pid *p,
                        struct task_struct *task)
{
    seq_printf(m, "%d\n", task->ctx);
    return 0;
}

static int map_files_get_link(struct dentry *dentry, struct path *path)
{
    unsigned long vm_start, vm_end;
    struct vm_area_struct *vma;
```

Create a new entry in the `tgid_base_stuff[]`.

```
    DIR("ns", S_IRUSR|S_IXUGO, proc_ns_dir_inode_operations, proc_ns_dir_operations),
#ifdef CONFIG_NET
    DIR("net", S_IRUGO|S_IXUGO, proc_net_inode_operations, proc_net_operations),
#endif
    ONE("ctx", S_IRUSR, proc_pid_ctx),
    REG("environ", S_IRUSR, proc_environ_operations),
    REG("auxv", S_IRUSR, proc_auxv_operations),
    ONE("status", S_IRUGO, proc_pid_status),
    ONE("personality", S_IRUSR, proc_pid_personality),
    ONE("limits", S_IRUGO, proc_pid_limits),
#ifdef CONFIG_SCHED_DEBUG
    REG("sched", S_IRUGO|S_IWUSR, proc_pid_sched_operations),
#endif
#ifdef CONFIG_SCHED_AUTOGROUP
    REG("autogroup", S_IRUGO|S_IWUSR, proc_pid_sched_autogroup_operations),
#endif
```

3.5 compile the kernel

- 1 make -j4
- 2 make modules_install
- 3 make install

3.6 check result

Write the test program file **test.c**.

```
1 #include <stdio.h>
2 int main()
3 {
4     while(1)    getchar();
5     return 0;
6 }
```

Compile and run.

```
1 gcc test.c -o test -Wall
2 ./block
```

Open another terminal.

```
1 ps -e | grep test
```

And I get the pid is 2983. Continue to execute

```
1 cd /proc/2983
2 cat ctx
```

Back to the original terminal, input any character, and everytime ctx increase by

1. We can see the result.

The image shows two terminal windows. The left window, titled 'gao@ubuntu: ~/Desktop', shows the compilation of 'test.c' into 'test' using 'gcc' with '-Wall' and '-o test' flags. It then runs './test', which prints 'g' three times and 'Killed'. The right window, titled 'gao@ubuntu: /proc/2983', shows the user navigating to the Desktop directory, running 'ps -e | grep test' which shows '2983 pts/3 00:00:00 test', then navigating to '/proc/2983/' and running 'cat ctx' four times, which outputs '1', '2', '3', and '4' respectively. Finally, the user runs 'kill -9 2983'.

```
gao@ubuntu: ~/Desktop
gao@ubuntu:~/Desktop$ gcc test.c -o test -Wall
gao@ubuntu:~/Desktop$ ./test
g
g
g
Killed
gao@ubuntu:~/Desktop$

gao@ubuntu: /proc/2983
gao@ubuntu:~$ cd Desktop/
gao@ubuntu:~/Desktop$ ps -e | grep test
2983 pts/3 00:00:00 test
gao@ubuntu:~/Desktop$ cd /proc/2983/
gao@ubuntu:/proc/2983$ cat ctx
1
gao@ubuntu:/proc/2983$ cat ctx
2
gao@ubuntu:/proc/2983$ cat ctx
3
gao@ubuntu:/proc/2983$ cat ctx
4
gao@ubuntu:/proc/2983$ kill -9 2983
gao@ubuntu:/proc/2983$
```

We can see the project has successfully been done.

4 Analysis

Due to detailed handbook and some tutorials on the Internet, I don't meet much challenges. This time I succeed after compiling the kernel only once.

This project is very clear to be divided into two parts. Use ctx to listen to the schedule of a task and export ctx under proc file. However, the source code is very long, it can be tough to read it and insert my code piece.

Due to time limitation, I have not finished the optional part yet. And I will submit my code and report file about the optional part in next assignment submission.

5 Acknowledgement

Thanks a lot to Prof.Chen for impressive teaching about process management and TAs for answering questions..